

YASH RAWAL

+91-9399159685 | yash.mahesh.rawal@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRL | Ratlam, MP, India

September 28, 2026

Hiring Team & Engineering Leadership
Applied AI & Machine Learning Engineering

Dear Hiring Team,

I am writing to express my strong interest in the **AI/ML Engineer / Applied AI Engineer** position. As an AI Engineer and Technical Lead with hands-on experience designing and operating production agentic platforms serving **1M+ active ERP users** across 120+ enterprise cloud tenants, I specialize in architecting reliable multi-agent systems, low-latency LLM serving infrastructure, and production RAG pipelines. Combining core model training with distributed systems engineering, my focus centers on building deterministic, secure, and cost-efficient AI applications that scale seamlessly in enterprise environments.

At Bigscal Technologies, I served as AI Lead with end-to-end technical ownership across architecture design, microservices, and model deployment. Key highlights of my technical work and leadership include:

- **Enterprise Multi-Agent Routing & Scale:** Architected multi-agent routing with LangGraph & SAP Joule 2.0 discovering and integrating 650+ dynamic MCP tools. Re-engineered execution flows with capability contracts and a unified ReAct loop, achieving **3× throughput gains and \$120K+ annual compute savings**.
- **High-Performance LLM Serving & Open-Source Contributions:** Active contributor to vLLM (merged PR #55475 fixing XML tool parser drops and resolved unhandled torchcodec multimodal video decoder crash boundaries). Fine-tuned and served Qwen2.5-7B and E5-large-v2 embeddings using PyTorch, Unsloth, and vLLM endpoints.
- **Hierarchical RAG & Zero-Trust Governance:** Built 4-level hierarchical RAG pipelines indexing 10,000+ technical documents via PGVector (achieving **90%+ Recall@10**), and engineered multi-tier zero-trust PII sanitization scrubbing sensitive data across tool and server boundaries before LLMs.
- **Evaluation Control Planes & Mentorship:** Conceived and developed the Prompt Lab control plane to capture execution traces and curate benchmark datasets for RLHF/DPO pipelines, while mentoring and upskilling 5+ junior engineers and interns in model fine-tuning and agent evaluation.

What distinguishes my engineering approach is the conviction that production AI requires systems-level rigor rather than brittle wrapper APIs. Whether pretraining sub-100M custom Transformers from scratch on 462M tokens, engineering chunked attention routines (FragmentStream) for memory-constrained hardware, or optimizing async inference request lifecycles, I focus on delivering tangible business value with strict performance and security guarantees.

I am excited about the opportunity to bring my hands-on architectural experience, model fine-tuning depth, and agent infrastructure leadership to your engineering team. Thank you for your time and consideration, and I welcome the opportunity to discuss how my background aligns with your engineering goals.

Sincerely,

Yash Rawal

AI/ML Engineer & Technical Lead