

Yash Rawal

+91-9399159685 | yashrawal987@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRL | Ratlam (457001), Madhya Pradesh, India

Professional Summary

Applied AI & Machine Learning Engineer serving a **1M+ user scale** by architecting enterprise agentic platforms, fine-tuning domain LLMs & embedding models (**Qwen2.5-7B, E5-large-v2**) on NVIDIA GPUs, and building automated evaluation control planes. Specialized in Python, high-throughput LLM serving internals (vLLM contributor), workflow orchestration (SAP Joule 2.0, LangGraph), and cloud-native infrastructure (Azure, AWS, Docker, Kubernetes).

Skills

Languages & Backend: Python, SQL, FastAPI, PostgreSQL, Oracle, REST APIs, WebSockets, AsyncIO, Linux

AI & ML Engineering: Agentic AI, Model Context Protocol (MCP), Multi-Agent Systems, Tool RAG, Healthcare NLP, Model Fine-Tuning, vLLM Serving

AI & Workflow Frameworks: LangChain, LangGraph, vLLM, PyTorch, Hugging Face Transformers, SAP Joule 2.0, LlamaIndex, GROBID

Cloud & Infrastructure: Azure, AWS, Docker, Kubernetes, CI/CD (GitHub Actions, Jenkins, Git, Gerrit), Redis, PGVector

Testing, QA & Security: PyTest, LLM-as-a-Judge, Prompt Lab, Red Teaming, PII Sanitization, Unit/Integration/Regression Testing

Work Experience

Bigscal Technologies Pvt. Ltd.

AI/ML Engineer

Oct 2024 - Present

Surat, India

- Deployed a production-scale agentic platform and multi-agent systems serving **1M+ active ERP users** across 120+ enterprise tenants on cloud infrastructure (Azure, AWS).
- Architected multi-agent routing with **LangGraph & SAP Joule 2.0**, implementing hybrid semantic-lexical tool discovery (vector + BM25F/RRF) across 650+ dynamic read/write and generation MCP capabilities.
- Designed a capability contract layer to execute nested profiles inside a single ReAct loop, reducing execution overhead by **3x** and cutting token consumption by **70%**, slashing enterprise inference costs from \$150 to \$30 per 1K multi-agent runs.
- Fine-tuned and served **Qwen2.5-7B & E5-large-v2** models on NVIDIA GPUs using PyTorch, Unsloth, and contrastive learning, deploying via vLLM for structured role mapping and semantic retrieval.
- Implemented a 4-level hierarchical RAG pipeline in Python indexing 10,000+ technical documents using GROBID, pdfplumber, and PostgreSQL/PGVector, achieving a **90%+ Recall@10** retrieval rate.
- Developed the Prompt Lab evaluation control plane in Python to collect reasoning trajectories and human preference annotations, establishing a scalable **Human Feedback (RLHF) Pipeline**.

Bigscal Technologies Pvt. Ltd.

AI/ML Engineer Intern

Apr 2024 - Sep 2024

Surat, India

- Engineered clinical summarization pipelines in Python for 116K+ medical records with 4-bit quantized Llama 3, and designed real-time WebSocket audio streaming for Indic ASR with **NVIDIA NeMo (RNN-T and CTC decoding)**.
- Researched audio digital signal processing (STFT, Mel-spectrograms, MFCC feature extraction via librosa/torchaudio) and CNN audio architectures (VGGish, ResNet), compiling custom Indic LJSpeech-format datasets.

Projects

vLLM High-Performance Inference Engine (vllm-project/vllm) | Open Source

Merged **PR #55475** fixing structural XML tool parser drops on `/v1/chat/completions` and Pydantic parameter indexing (64/64 Qwen tests passed); engineered `'torchcodec'` multimodal decoder guards (Issue #54097).

Python, PyTorch, Pydantic, AsyncIO, PyTest

Job Architecture Generator (Qwen2.5-7B Domain SFT & Serving) | LLM Training

Fine-tuned **Qwen2.5-7B** on 10K+ domain examples using Unsloth & PyTorch on NVIDIA GPUs for role standardization and structured JSON generation; deployed an **OpenAI-compatible endpoint** via vLLM for low-latency enterprise inference.

Python, PyTorch, Hugging Face Transformers, Unsloth, vLLM, NVIDIA GPUs

E5-large-v2 Job & Skills Embedding Fine-Tuning | Model Training

Fine-tuned **E5-large-v2** (335M parameters) on Lightcast labor-market data using PyTorch and triplet loss (anchor + positive + hard negative), raising semantic retrieval relevance by **18%** over OpenAI text-embedding-3-small.

Python, PyTorch, Hugging Face Transformers, Triplet Loss, NVIDIA GPUs

Modular Agent Harness | Project

Principal Creator & Architect of a modular Python agentic framework implementing a formal ReAct loop with step budgets, pgvector tool routing, and dynamic MCP settings caching.

Python, FastAPI, PostgreSQL, WebSockets, pgvector

Enterprise Hierarchical Document Intelligence & RAG Platform | Project

Built an immutable document versioning system and 4-level hierarchical ingestion pipeline (paragraphs, pages, chapters, headings) using GROBID and pdfplumber with paragraph-level signatures in `'pgvector'` (**90%+ Recall@10**).

Python, PostgreSQL, pgvector, GROBID, pdfplumber

Education

Master of Computer Applications (Artificial Intelligence & Machine Learning) - Parul University

2022 - 2024

Certifications

Microsoft Certified: Azure AI Fundamentals (AI-900)

Microsoft

Microsoft Certified: Azure Fundamentals (AZ-900)

Microsoft