

YASH RAWAL

+91-9399159685 | yash.mahesh.rawal@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRI | Ratlam, MP, India

PROFESSIONAL SUMMARY

AI/ML Engineer specializing in production LLMs, Agentic AI architectures, and low-latency RAG systems with PyTorch and vLLM. Proven expertise in fine-tuning open-source models, building multi-agent MCP tools, and deploying enterprise AI pipelines.

EXPERIENCE

Bigscal Technologies Pvt. Ltd.

Surat, India

AI/ML Engineer

Oct. 2024 – Present

- Deployed enterprise agentic platform serving 1M+ active ERP users across 120+ cloud tenants on Azure and AWS.
- Architected multi-agent routing with LangGraph & SAP Joule 2.0, discovering and integrating 650+ dynamic MCP tools.
- Re-architected multi-agent execution with capability contracts, yielding 3× throughput gains and \$120K+ annual savings.
- Fine-tuned and served Qwen2.5-7B and E5-large-v2 on NVIDIA GPUs using PyTorch, Unsloth, and vLLM inference endpoints.
- Built 4-level hierarchical RAG in Python indexing 10,000+ technical documents via PGVector achieving 90%+ Recall@10.
- Built Prompt Lab evaluation control plane to monitor LLM outputs, track latency, and curate RLHF benchmark datasets.

Bigscal Technologies Pvt. Ltd.

Surat, India

AI/ML Intern

Apr. 2024 – Sep. 2024

- Engineered clinical NLP summarization pipelines for 116K+ medical records using 4-bit quantized Llama 3 and FastAPI.
- Built real-time WebSocket audio streaming for Indic ASR using NVIDIA NeMo (RNN-T and CTC decoding) and audio DSP.

PROJECTS

vLLM High-Performance Inference Engine | Python, PyTorch, Pydantic, AsyncIO

Open Source

- Merged PR #55475 fixing XML tool parser drops on /v1/chat/completions and engineered torchcodec compatibility guards.

Job Architecture Generator (Qwen2.5-7B SFT) | Python, PyTorch, Unsloth, vLLM

LLM Training

- Fine-tuned Qwen2.5-7B on 10K+ domain examples using Unsloth and deployed low-latency OpenAI-compatible endpoints.

E5-large-v2 Job & Skills Embedding Fine-Tuning | Python, PyTorch, Transformers

Model Training

- Fine-tuned E5-large-v2 (335M params) with triplet loss, boosting domain semantic retrieval relevance by 18% over OpenAI.

Modular Agent Harness & Conversational Copilot | Python, FastAPI, pgvector

Project

- Architected modular Python agent framework implementing formal ReAct loops with step budgets and dynamic tool routing.

Hierarchical Document Intelligence & RAG Platform | Python, pgvector, GROBID

Project

- Engineered 4-level hierarchical RAG ingestion for complex documents using GROBID parsing and PGVector similarity search.

TECHNICAL SKILLS

Languages & Backend: Python, SQL, FastAPI, PostgreSQL, Oracle, REST APIs, WebSockets, AsyncIO, Linux, Git, Docker

AI & ML Engineering: Large Language Models (LLMs), Agentic AI, MCP, Tool RAG, SFT/LoRA, vLLM Serving, RLHF, Prompt Tuning

Frameworks & Libraries: LangChain, LangGraph, PyTorch, Hugging Face Transformers, vLLM, SAP Joule 2.0, NVIDIA NeMo

Cloud & DevOps: Azure, AWS, Docker, Kubernetes, CI/CD (GitHub Actions), Redis, PGVector, PyTest, LLM-as-a-Judge

EDUCATION & CERTIFICATIONS

Parul University

Gujarat, India

Master of Computer Applications (Artificial Intelligence & Machine Learning)

2022 – 2024

Certifications: Microsoft Certified: Azure AI Fundamentals (AI-900) | Microsoft Certified: Azure Fundamentals (AZ-900)