

Yash Rawal

+91-9399159685 | yashrawal987@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRL | Surat, Gujarat, India

Professional Summary

AI Platform Engineer architecting enterprise LLM infrastructure and multi-agent systems for 1M+ active users. Specialized in building low-latency, scalable agent execution runtimes, zero-vector-DB tool-routing layers, and hierarchical RAG pipelines. Open-source contributor to Spring AI, leveraging deep research experience and system design to ship production-scale software.

Skills

Languages & Backend: Python, Java, SQL, Spring Boot, FastAPI, PostgreSQL, SAP HANA
AI Frameworks & Platforms: LangChain, LangGraph, LlamaIndex, Spring AI, CrewAI, AutoGen, Docker, Kubernetes
Evaluation & Security: LLM-as-Judge, Prompt Lab, Red Teaming, Prompt Injection Guardrails, RLHF, DPO, Golden Datasets

Work Experience

TalenTeam Ltd., UK (via Bigscal Technologies)

AI/ML Engineer Oct 2024 - Present

Remote

- Architected and deployed production-scale applied AI workflows and multi-agent family graphs (flat orchestration and Model Context Protocol tools) serving 1M+ active ERP users across 120+ enterprise organizations.
- Addressed Ask Engine multi-agent latency by designing a capability contract layer that executes nested capability profiles inside a single ReAct loop, reducing latency by 3x and token costs by 70%.
- Built a custom zero-vector-DB tool routing system on volatile RAM (using BM25F, TF-IDF, Elias encoding, and LSA fallback) to scale 1,000+ capabilities with microsecond-level latency.
- Architected the Prompt Lab evaluation system for large-scale model evaluations by engineering a FastAPI + UI + n8n control plane to capture durable trajectories and human preference annotations, serving as an offline RLHF/DPO dataset factory to optimize routing policies.
- Built a search platform for 10,000+ documents at an international patent authority by implementing a hierarchical RAG pipeline (indexing paragraph, page, and chapter layers) and Postgres signature deduplication, achieving 90%+ Recall@10.
- Created enterprise agent workflows for frontline users by conducting user research and usability testing to integrate agents with ERP, CRM, and scheduling systems, shipping software with high adoption.

Healthray (Bigscal Technologies)

AI/ML Engineer Intern Apr 2024 - Sep 2024

Surat, India

- Built multilingual clinical workflows (Whisper transcription, medical NER, speech synthesis) with rule-based safety guardrails.
- Developed LLM-based safeguards for structured extraction of medications and patient history.

Projects

Spring AI (spring-projects/spring-ai) | Active Contributor

Merged DeepSeek-R1/Qwen reasoningContent metadata fix (PR #5711) and prompt-cache TTL support for AWS Bedrock Converse API to resolve streaming/non-streaming consistency.

Java, Spring Boot, AWS Bedrock

Model Context Protocol (MCP) & Claude Code | Research & Integration

Researched Model Context Protocol (MCP) tool routing integrations with Claude Code, Cursor, and custom agentic workflows to secure API token validation.

Claude Code, Cursor, JSON-RPC 2.0

AI Agent Security Red Teaming (OpenAI Competition) | Research

Developed automated adversarial pipelines (bypassing guardrails and security observers via optimized multi-predicate prompt injections) to achieve level 5 severity exploits across 2,000+ test cases.

Kaggle, OpenAI API

Education

Master of Computer Applications (MCA - AI & ML) - Parul University 2024

Certifications

Microsoft Certified: Azure AI Fundamentals (AI-900)	Microsoft	2024
Microsoft Certified: Azure Fundamentals (AZ-900)	Microsoft	2024