

Yash Rawal

Ratlam (457001), Madhya Pradesh, India | yashrawal987@gmail.com | +91-9399159685 | linkedin.com/in/rawal-yash | github.com/YashRL

Master of Computer Applications (AI & ML) | Microsoft Certified: Azure AI Fundamentals (AI-900)

Cover Letter: GenAI Engineer (Job ID: 1970324837507805)

Ascendion Pune, India September 2, 2026

Dear Ascendion Hiring Team,

I am writing to express my enthusiastic interest in the GenAI Engineer role at Ascendion in Pune (Job ID: 1970324837507805).

As an Applied AI & Machine Learning Engineer currently architecting enterprise agentic platforms and automated intelligence systems serving 1M+ active users across 120+ organizations, my technical work directly aligns with Ascendion's AI-native engineering ethos and your 'Engineering to the Power of AI' methodology. Combining deep hands-on expertise in LangChain and LangGraph orchestration, hierarchical RAG architectures, PostgreSQL/PGVector, and model evaluation control planes, I specialize in translating complex business requirements into scalable, secure, and production-ready Generative AI solutions.

- **Autonomous Agentic AI & ReAct Orchestration**, Architected production-scale multi-agent workflows and designed a single ReAct loop capability contract layer that cut execution latency by 3x and token consumption by 70% across enterprise tenants.
- **Hierarchical RAG & Vector Knowledge Retrieval**, Implemented a hierarchical RAG pipeline indexing 10,000+ technical documents using GROBID, pdfplumber, and PostgreSQL/PGVector, achieving a 90%+ Recall@10 retrieval rate.
- **Model Evaluation & Quality Governance**, Developed the Prompt Lab evaluation control plane to systematically log multi-turn reasoning trajectories, tool call payloads, and human preference annotations, establishing a scalable Human Feedback (RLHF) Pipeline.
- **Model Fine-Tuning & Open-Source Depth**, Trained domain-specific Qwen2.5-7B models and fine-tuned dense embeddings using PyTorch; contributed upstream to Spring AI (PR #5711) resolving streaming reasoningContent and Converse API prompt-caching latency.

Delivering measurable business outcomes for Global 2000 enterprises through thousands of AI agents requires robust software engineering discipline, deterministic tool execution, and rigorous evaluation pipelines. Having deployed high-concurrency agentic runtimes at 1M+ user scale and fortified them against OWASP LLM vulnerabilities with runtime tool monitoring and automated PII redaction, I am eager to help Ascendion engineer the next generation of AI-native platform solutions.

Thank you for your time and consideration. I would welcome the opportunity to discuss how my technical background in Generative AI engineering, RAG optimization, and agentic workflows can add immediate value to Ascendion's engineering team.

Sincerely,
Yash Rawal