

Yash Rawal

+91-9399159685 | yashrawal987@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRL | Surat, Gujarat, India

Professional Summary

AI Engineer specialized in AI Platform Engineering and enterprise LLM infrastructure for 1M+ active users. Skilled in building low-latency, secure agent execution runtimes, zero-vector-DB tool-routing layers, and AWS-native AI services. Open-source contributor to Spring AI, leveraging deep system design and security experience to ship production-scale software.

Skills

Languages & Backend: Python, Java, SQL, Spring Boot, FastAPI, PostgreSQL, SAP HANA
AI Frameworks & Platforms: LangChain, LangGraph, LlamaIndex, Spring AI, CrewAI, AutoGen, Docker, Kubernetes
Cloud & Security: AWS Bedrock, AWS SageMaker, Azure OpenAI, CI/CD, LLM-as-Judge, OWASP LLM Top 10, PII Sanitization, Red Teaming

Work Experience

TalenTeam Ltd., UK (via Bigscal Technologies)

AI/ML Engineer Oct 2024 - Present

- Remote
- Developed full-stack applied AI workflows and multi-agent family graphs (flat orchestration and Model Context Protocol tools) on AWS Bedrock serving 1M+ active ERP users across healthcare, manufacturing, and financial services.
 - Addressed Ask Engine multi-agent latency by designing a capability contract layer that executes nested capability profiles inside a single ReAct loop, reducing latency by 3x and token costs by 70%.
 - Built a custom zero-vector-DB tool routing system on volatile RAM (using BM25F, TF-IDF, Elias encoding, and LSA fallback) to scale 1,000+ capabilities with microsecond-level latency.
 - Architected the Prompt Lab evaluation system for large-scale model evaluations by engineering a FastAPI + UI + n8n control plane to capture durable trajectories and human preference annotations, serving as an offline RLHF/DPO dataset factory to optimize routing policies.
 - Built a search platform for 10,000+ documents at an international patent authority by implementing a hierarchical RAG pipeline (indexing paragraph, page, and chapter layers) and Postgres signature deduplication, achieving 90%+ Recall@10.
 - Secured agentic workflows against OWASP LLM Top 10 vulnerabilities by implementing observer-based tool-call monitoring and automated PII redaction, protecting sensitive customer financial data in production.

Healthray (Bigscal Technologies)

AI/ML Engineer Intern Apr 2024 - Sep 2024

- Surat, India
- Built clinical transcription workflows (Whisper transcription, medical NER, speech synthesis) using rule-based guardrails to extract medication schedules.
 - Developed LLM-based safeguards for structured extraction of medications and patient history.

Projects

Spring AI (spring-projects/spring-ai) | Active Contributor

Merged DeepSeek-R1/Qwen reasoningContent metadata fix (PR #5711) and prompt-cache TTL support for AWS Bedrock Converse API to resolve streaming/non-streaming consistency.
Java, Spring Boot, AWS Bedrock

Model Context Protocol (MCP) & Claude Code | Research & Integration

Researched Model Context Protocol (MCP) tool routing integrations with Claude Code, Cursor, and custom agentic workflows to secure API token validation.
Claude Code, Cursor, JSON-RPC 2.0

AI Agent Security Red Teaming (OpenAI Competition) | Research

Developed automated adversarial pipelines (bypassing guardrails and security observers via optimized multi-predicate prompt injections) to achieve level 5 severity exploits across 2,000+ test cases.
Kaggle, OpenAI API

Education

Master of Computer Applications (MCA - AI & ML) - Parul University 2024

Certifications

Microsoft Certified: Azure AI Fundamentals (AI-900)	Microsoft	2024
Microsoft Certified: Azure Fundamentals (AZ-900)	Microsoft	2024