

# Yash Rawal

Ratlam (457001), Madhya Pradesh, India | [yashrawal987@gmail.com](mailto:yashrawal987@gmail.com) | +91-9399159685 | [linkedin.com/in/rawal-yash](https://linkedin.com/in/rawal-yash) | [github.com/YashRL](https://github.com/YashRL)

Master of Computer Applications (AI & ML) | Microsoft Certified: Azure AI Fundamentals (AI-900)

---

## Cover Letter: Staff AI Engineer — Servicing Enablement Platform (Job ID: 26011938)

American Express Bengaluru, India September 1, 2026

Dear American Express Hiring Team,

I am writing to express my strong interest in the Staff AI Engineer position within the Service Enablement Platform (SEP) team at American Express (Job ID: 26011938).

As an Applied AI & Machine Learning Engineer currently architecting enterprise agentic platforms and multi-agent systems serving 1M+ active users across 120+ organizations, my technical focus directly aligns with Amex's vision of deploying autonomous agentic systems that plan, reason, and act across mission-critical financial workflows. Combining deep hands-on expertise in LangGraph orchestration, hierarchical RAG architectures, model fine-tuning, and robust evaluation control planes, I specialize in building reliable, low-latency, and compliant agentic frameworks for enterprise scale.

- **Enterprise Agentic Platform Architecture**, Architected production-scale multi-agent workflows and engineered a single ReAct loop capability contract layer that cut execution latency by 3x and token consumption by 70% across enterprise tenants.
- **Hierarchical RAG & Enterprise Grounding**, Implemented a hierarchical RAG pipeline indexing 10,000+ technical documents using GROBID, pdfplumber, and PostgreSQL/PGVector, achieving a 90%+ Recall@10 retrieval rate.
- **Evaluation, Safety & Observability Control Planes**, Developed the Prompt Lab evaluation control plane to systematically log multi-turn reasoning trajectories, tool call payloads, and human preference annotations, establishing a scalable Human Feedback (RLHF) Pipeline.
- **Financial-Grade Security & Governance**, Fortified agentic workflows against OWASP LLM vulnerabilities using phantom-token context propagation, runtime tool monitoring, and trained token-classification models for automated PII redaction.

Deploying agentic AI within the Service Enablement Platform across disputes management, fraud servicing, and credit operations requires more than basic model calls—it demands deterministic tool execution, strict autonomy boundaries, rigorous PII protection, and auditable reasoning trajectories that meet financial regulatory standards. Having architected and operated high-throughput agentic platforms at 1M+ user scale and contributed upstream to Spring AI, I am eager to help Amex define and scale its next-generation agentic AI capabilities.

Thank you for your time and consideration. I would welcome the opportunity to discuss how my hands-on architectural leadership in agentic systems, RAG optimization, and enterprise AI safety can accelerate the mission of the Service Enablement Platform team at American Express.

Sincerely,  
Yash Rawal