

YASH RAWAL

+91-9399159685 | yash.mahesh.rawal@gmail.com | [linkedin.com/in/rawal-yash](https://www.linkedin.com/in/rawal-yash) | github.com/YashRL | Ratlam, MP, India

PROFESSIONAL SUMMARY

AI Engineer & Technical Lead specializing in enterprise GenAI, Agentic AI architectures, and low-latency RAG systems. Proven expertise in Azure cloud integrations, LLMOps pipelines, Responsible AI guardrails, and cost-efficient scale.

EXPERIENCE

Bigscal Technologies Pvt. Ltd.

Surat, India

AI/ML Engineer (Technical Lead)

Oct. 2024 – Present

- Served as AI Lead driving 100% end-to-end ownership across multi-tenant microservices, agent runtimes, and models.
- Deployed enterprise agentic platform serving 1M+ active ERP users across 120+ cloud tenants on Azure and AWS.
- Architected multi-agent A2A routing with LangGraph & SAP Joule 2.0, unifying SAP SuccessFactors, Blend, and MCP APIs.
- Engineered dynamic UI widget rendering contracts and capability layers, boosting throughput 3× and saving \$120K+/year.
- Fine-tuned and served Qwen2.5-7B and E5-large-v2 on NVIDIA GPUs using PyTorch, Unsloth, and vLLM inference endpoints.
- Built 4-level hierarchical RAG in Python indexing 10,000+ technical documents via PGVector achieving 90%+ Recall@10.
- Conceived and built Prompt Lab evaluation control plane to benchmark LLM outputs and track hallucination metrics.
- Architected multi-tier zero-trust PII sanitization and safety guardrails scrubbing sensitive data before LLMs.
- Mentored and upskilled 5+ interns and junior engineers in PyTorch fine-tuning, agentic workflows, and LLM evaluation.

Bigscal Technologies Pvt. Ltd.

Surat, India

AI/ML Intern

Apr. 2024 – Sep. 2024

- Engineered clinical NLP summarization pipelines for 116K+ medical records using 4-bit quantized Llama 3 and FastAPI.
- Built real-time WebSocket audio streaming for Indic ASR using NVIDIA NeMo (RNN-T and CTC decoding) and audio DSP.
- Constructed multilingual speech synthesis pipelines with Coqui TTS, formatting custom LJSpeech audio datasets via DSP.

PROJECTS

vLLM High-Performance Inference Engine | Python, PyTorch, Pydantic, AsyncIO

Open Source

- Merged PR #55475 fixing XML tool parser drops on /v1/chat/completions and engineered torchcodec compatibility guards.
- Resolved unhandled torchcodec multimodal video decoder crashes by engineering resilient FFmpeg error boundaries.
- Optimized async inference request lifecycle and validated Pydantic schema serialization across multi-GPU test suites.

Job Architecture Generator (Qwen2.5-7B SFT) | Python, PyTorch, Unsloth, vLLM

LLM Training

- Fine-tuned Qwen2.5-7B on 10K+ domain examples using Unsloth and deployed low-latency OpenAI-compatible endpoints.
- Trained 4-bit QLoRA adapters with custom chat templates, cutting GPU memory overhead by 60% during vLLM serving.

E5-large-v2 Job & Skills Embedding Fine-Tuning | Python, PyTorch, Transformers

Model Training

- Fine-tuned E5-large-v2 (335M params) with triplet loss, boosting domain semantic retrieval relevance by 18% over OpenAI.
- Constructed hard-negative mining pipeline over Lightcast labor datasets to separate closely aligned skill taxonomies.

Edge Language Model for Edge Hardware | Python, PyTorch, CUDA, TinyStories

Research Project

- Pretrained sub-100M autoregressive Transformer from scratch on 462M tokens with custom Pre-LN architecture (1.43 val loss).
- Engineered FragmentStream tiled memory attention algorithm preventing OOMs during long-context generation.

Modular Agent Harness & Conversational Copilot | Python, FastAPI, pgvector

Project

- Architected modular Python agent framework implementing formal ReAct loops with step budgets and dynamic tool routing.
- Engineered persistent multi-tenant session state and automated checkpoint recovery across asynchronous worker threads.

Hierarchical Document Intelligence & RAG Platform | Python, pgvector, GROBID

Project

- Engineered 4-level hierarchical RAG ingestion for complex documents using GROBID parsing and PGVector similarity search.
- Developed paragraph-level signature deduplication and versioning pipeline supporting 1,000+ domain reference books.

TECHNICAL SKILLS

Languages & Backend: Python, SQL, Spark, FastAPI, PostgreSQL, Oracle, REST APIs, WebSockets, AsyncIO, Linux, Git, Docker

AI & ML Engineering: Generative AI, LLMs, Agentic AI, MCP, RAG, Responsible AI, Safety Guardrails, SFT/LoRA, vLLM Serving

Frameworks & Libraries: LangChain, LangGraph, PyTorch, Hugging Face Transformers, vLLM, SAP Joule 2.0, NVIDIA NeMo

Cloud & MLOps/LLMOps: Azure, AWS, Databricks, Kubernetes, CI/CD (GitHub Actions), Redis, PGVector, PyTest, LLM-as-a-Judge

EDUCATION & CERTIFICATIONS

Parul University

Gujarat, India

Master of Computer Applications (Artificial Intelligence & Machine Learning)

2022 – 2024

Relevant Coursework: Deep Learning, Natural Language Processing, Distributed Systems, Cloud Architecture, Algorithms

Certifications: Microsoft Certified: Azure AI Fundamentals (AI-900) | Microsoft Certified: Azure Fundamentals (AZ-900)