

Yash Rawal

+91-9399159685 | yashrawal987@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRL | Ratlam (457001), Madhya Pradesh, India

Professional Summary

Applied AI & Machine Learning Engineer serving a **1M+ user scale** by architecting enterprise agentic platforms, hierarchical RAG architectures, and automated intelligence pipelines across 120+ organizations. Specialized in Python, Model Context Protocol (MCP), multi-agentic orchestration (LangChain, LangGraph), and healthcare NLP. Active contributor to Spring AI, bridging generative AI and model fine-tuning with resilient cloud deployments (Azure, AWS).

Skills

Languages & Backend: Python, SQL, PostgreSQL, Java, FastAPI, Spring Boot, REST APIs

AI & ML Engineering: Agentic AI, Model Context Protocol (MCP), Multi-Agent Systems, Tool RAG, Healthcare NLP

AI & ML Frameworks: LangChain, LangGraph, Spring AI, PyTorch, vLLM, LlamaIndex, GROBID

Platforms & Models: Azure, AWS, Azure OpenAI, Amazon Bedrock, Meta Llama, Claude, Qwen

Infrastructure & Security: Docker, Kubernetes, Linux, CI/CD, LLM-as-Judge, Red Teaming, PII Sanitization

Work Experience

Bigscal Technologies Pvt. Ltd.

AI/ML Engineer

Oct 2024 - Present

Surat, India

- Deployed a production-scale agentic platform and multi-agent systems serving **1M+ active ERP users** across 120+ enterprise tenants on cloud infrastructure (Azure, AWS).
- Built production guardrails against prompt injections and data leaks, integrating runtime tool validation and **automated PII redaction**.
- Designed a capability contract layer to execute nested profiles inside a single ReAct loop, reducing execution overhead by **3x** and token consumption by 70%.
- Implemented a hierarchical RAG pipeline indexing 10,000+ technical documents using GROBID, pdfplumber, and PostgreSQL/PGVector, achieving a **90%+ Recall@10** retrieval rate.
- Engineered an automated LLM evaluation and red-teaming control plane to log reasoning trajectories, scoring agent behaviors via LLM-as-a-judge and **RLHF preference datasets**.
- Built production guardrails against prompt injections and data leaks, integrating runtime tool validation and **automated PII redaction**.

Bigscal Technologies Pvt. Ltd.

AI/ML Engineer Intern

Apr 2024 - Sep 2024

Surat, India

- Engineered clinical summarization pipelines for 116K+ medical records with 4-bit quantized Llama 3, and designed real-time WebSocket audio streaming for Indic ASR with **NeMo RNN-T and CTC decoding**.
- Researched audio digital signal processing (STFT, Mel-spectrograms, MFCC feature extraction via librosa/torchaudio) and CNN audio architectures (VGGish, ResNet), compiling custom Indic LJSpeech-format datasets.

Projects

Modular Agent Harness & Memory System ("Brain") | Open Source

Principal Creator & Architect of an open-source plug-and-play agentic framework designed as modular building blocks for flexible enterprise workflows, implementing a formal ReAct loop with step budgets, vector tool routing, and dynamic MCP settings caching.

Java, Spring Boot, PostgreSQL, WebSockets

Spring AI (spring-projects/spring-ai) | Open Source

Merged PR #5711 to propagate reasoningContent from AWS Bedrock and OpenAI-compatible models, and resolved Converse API prompt-caching TTL latency consistency.

Java, Spring Boot, Python

Mini Language Model (~85M Parameter Architecture) | Personal Research

Pretrained an ~85M-parameter transformer to ~1.43 validation loss on TinyStories targeting **edge hardware deployment**; engineered FragmentStream chunked attention to eliminate $O(T^2)$ VRAM overhead and custom Triton kernels with KV caching for low-latency inference.

Python, PyTorch, CUDA, Triton, Tesla P100/T4

Education

Master of Computer Applications (Artificial Intelligence & Machine Learning) - Parul University

2022 - 2024

Certifications

Microsoft Certified: Azure AI Fundamentals (AI-900)

Microsoft

Microsoft Certified: Azure Fundamentals (AZ-900)

Microsoft