

```
{ "lang": "en", "page_format": "a4", "name": "Yash Rawal", "contact_line": "Ratlam (457001),  
Madhya Pradesh, India | +91-9399159685", "email": { "url": "mailto:yashrawal987@gmail.com",  
"display": "yashrawal987@gmail.com" }, "linkedin": { "url": "https://linkedin.com/in/rawal-yash",  
"display": "linkedin.com/in/rawal-yash" }, "github": { "url": "https://github.com/YashRL", "display":  
"github.com/YashRL" }, "candidate": { "name": "Yash Rawal", "phone": "+91-9399159685", "email":  
"yashrawal987@gmail.com", "linkedin": { "url": "https://linkedin.com/in/rawal-yash", "display":  
"linkedin.com/in/rawal-yash" }, "github": { "url": "https://github.com/YashRL", "display":  
"github.com/YashRL" }, "portfolio": null, "location": "Ratlam (457001), Madhya Pradesh, India" },  
"summary": "Applied AI & Machine Learning Engineer serving a 1M+ user scale by architecting  
enterprise agentic platforms, hierarchical RAG architectures, and automated intelligence pipelines  
across 120+ organizations. Specialized in Python, workflow orchestration (n8n, LangGraph), Model  
Context Protocol (MCP), and automated testing/evaluation control planes. Active contributor to Spring  
AI, bridging generative AI and model fine-tuning with resilient cloud deployments (Azure, AWS).",  
"experience": [ { "company": "Bigscal Technologies Pvt. Ltd.", "role": "AI/ML Engineer", "location":  
"Surat, India", "dates": "Oct 2024 - Present", "bullets": [ "Deployed a production-scale agentic  
platform and multi-agent systems serving 1M+ active ERP users across 120+ enterprise tenants on  
cloud infrastructure (Azure, AWS).", "Architected multi-agent routing and workflow automation with  
LangGraph & n8n, implementing hybrid semantic-lexical tool discovery (vector + BM25F/RRF)  
across 65+ MCP capabilities.", "Designed a capability contract layer to execute nested profiles inside a  
single ReAct loop, reducing execution overhead by 3x and token consumption by 70%.",  
"Implemented a hierarchical RAG pipeline indexing 10,000+ technical documents using GROBID,  
pdfplumber, and PostgreSQL/PGVector, achieving a 90%+ Recall@10 retrieval rate.", "Developed the  
Prompt Lab evaluation control plane to collect reasoning trajectories and human preference  
annotations, establishing a scalable Human Feedback (RLHF) Pipeline.", "Built production  
guardrails against prompt injections and data leaks, integrating runtime tool validation and automated  
PII redaction." ] }, { "company": "Bigscal Technologies Pvt. Ltd.", "role": "AI/ML Engineer Intern",  
"location": "Surat, India", "dates": "Apr 2024 - Sep 2024", "bullets": [ "Engineered clinical  
summarization pipelines for 116K+ medical records with 4-bit quantized Llama 3, and designed real-  
time WebSocket audio streaming for Indic ASR with NeMo RNN-T and CTC decoding.",  
"Researched audio digital signal processing (STFT, Mel-spectrograms, MFCC feature extraction via  
librosa/torchaudio) and CNN audio architectures (VGGish, ResNet), compiling custom Indic  
LJSpeech-format datasets." ] } ], "projects": [ { "name": "Modular Agent Harness & Memory System  
('Brain')", "badge": "Open Source", "tech": "Java, Spring Boot, PostgreSQL, WebSockets",  
"description": "Principal Creator & Architect of an open-source plug-and-play agentic framework  
designed as modular building blocks for flexible enterprise workflows, implementing a formal ReAct  
loop with step budgets, vector tool routing, and dynamic MCP settings caching." }, { "name": "Spring  
AI (spring-projects/spring-ai)", "badge": "Open Source", "tech": "Java, Spring Boot, Python",  
"description": "Merged PR #5711 to propagate reasoningContent from AWS Bedrock and OpenAI-  
compatible models, and resolved Converse API prompt-caching TTL latency consistency." }, {  
"name": "Mini Language Model (~85M Parameter Architecture)", "badge": "Personal Research",  
"tech": "Python, PyTorch, CUDA, Triton, Tesla P100/T4", "description": "Pretrained an ~85M-  
parameter transformer to ~1.43 validation loss on TinyStories targeting edge hardware deployment;  
engineered FragmentStream chunked attention to eliminate O(T^2) VRAM overhead and custom  
Triton kernels with KV caching for low-latency inference." } ], "education": [ { "title": "Master of  
Computer Applications (Artificial Intelligence & Machine Learning)", "org": "Parul University",  
"year": "2022 - 2024" } ], "certifications": [ { "title": "Microsoft Certified: Azure AI Fundamentals  
(AI-900)", "org": "Microsoft" }, { "title": "Microsoft Certified: Azure Fundamentals (AZ-900)", "org":  
"Microsoft" } ], "skills": [ { "category": "Languages & Backend", "items": "Python, SQL, PostgreSQL,  
Java, FastAPI, Spring Boot, REST APIs, Linux" }, { "category": "AI & ML Engineering", "items":  
"Agentic AI, Model Context Protocol (MCP), Multi-Agent Systems, Tool RAG, Healthcare NLP" }, {  
"category": "AI & Workflow Frameworks", "items": "LangChain, LangGraph, n8n, Spring AI,  
PyTorch, vLLM, LlamaIndex, GROBID" }, { "category": "Cloud & Infrastructure", "items": "Azure,  
AWS, Docker, Kubernetes, CI/CD (Git/GitHub/Jenkins), Redis, PGVector" }, { "category": "Testing &  
Governance", "items": "LLM-as-a-Judge, Prompt Lab, Red Teaming, PII Sanitization, Unit/Integration  
Testing" } ] }
```