

Yash Rawal

+91-9399159685 | yashrawal987@gmail.com | linkedin.com/in/rawal-yash | github.com/YashRL | Surat, Gujarat, India

Professional Summary

Generative AI Engineer specialized in AI Platform Engineering and enterprise LLM infrastructure for 1M+ active users. Skilled in building low-latency, secure agent execution runtimes, zero-vector-DB tool-routing layers, and automated evaluation pipelines. Open-source contributor to Spring AI, leveraging deep system design and security experience to ship production-scale software.

Skills

Languages & Backend: Python (Async), Java, SQL, FastAPI, Flask, PostgreSQL, SAP HANA, REST APIs, Data Engineering
AI & Systems Infrastructure: LangGraph, LangChain, CrewAI, RAG, pgvector, FAISS, Docker, Kubernetes, CI/CD, Observability, GDPR Security
Evaluation & Research: LLM-as-Judge, Red Teaming, Prompt Injection Defence, RLHF, DPO, Golden Datasets

Work Experience

TalenTeam Ltd., UK (via Bigscal Technologies)

AI/ML Engineer	Oct 2024 - Present
Remote	
- Developed full-stack applied AI workflows and multi-agent family graphs (flat orchestration and Model Context Protocol tools) to automate IT operations (AIOps) serving 1M+ active ERP users across telecommunications and digital services. - Addressed Ask Engine multi-agent latency by designing a capability contract layer that executes nested capability profiles inside a single ReAct loop, reducing latency by 3x and token costs by 70%. - Built a custom zero-vector-DB tool routing system on volatile RAM (using BM25F, TF-IDF, Elias encoding, and LSA fallback) to scale 1,000+ capabilities with microsecond-level latency. - Architected the Prompt Lab evaluation system for large-scale model evaluations by engineering a FastAPI + UI + n8n control plane to capture durable trajectories and human preference annotations, establishing observability and serving as an offline RLHF/DPO dataset factory to optimize routing policies. - Built a search platform for 10,000+ documents at an international patent authority by implementing a hierarchical RAG pipeline (indexing paragraph, page, and chapter layers) and Postgres signature deduplication, achieving 90%+ Recall@10. - Secured agentic workflows against OWASP LLM Top 10 vulnerabilities by implementing observer-based tool-call monitoring and automated PII redaction, protecting data privacy and sensitive customer financial details.	

Healthray (Bigscal Technologies)

AI/ML Engineer Intern	Apr 2024 - Sep 2024
Surat, India	
- Built clinical transcription workflows (Whisper transcription, medical NER, speech synthesis) using rule-based guardrails to extract medication schedules. - Developed LLM-based safeguards for structured extraction of medications and patient history.	

Projects

Spring AI (spring-projects/spring-ai) Active Contributor	
Merged DeepSeek-R1/Qwen reasoningContent metadata fix (PR #5711) and prompt-cache TTL support for AWS Bedrock Converse API to resolve streaming/non-streaming consistency.	
Java, Spring Boot, AWS Bedrock	
Model Context Protocol (MCP) & Claude Code Research & Integration	
Researched Model Context Protocol (MCP) tool routing integrations with Claude Code, Cursor, and custom agentic copilots to secure API token validation.	
Claude Code, Cursor, JSON-RPC 2.0	
AI Agent Security Red Teaming (OpenAI Competition) Research	
Developed automated adversarial pipelines (bypassing guardrails and security observers via optimized multi-predicate prompt injections) to achieve level 5 severity exploits across 2,000+ test cases.	
Kaggle, OpenAI API	

Education

Master of Computer Applications (MCA - AI & ML) - Parul University	2024
---	------

Certifications

Microsoft Certified: Azure AI Fundamentals (AI-900)	Microsoft	2024
Microsoft Certified: Azure Fundamentals (AZ-900)	Microsoft	2024